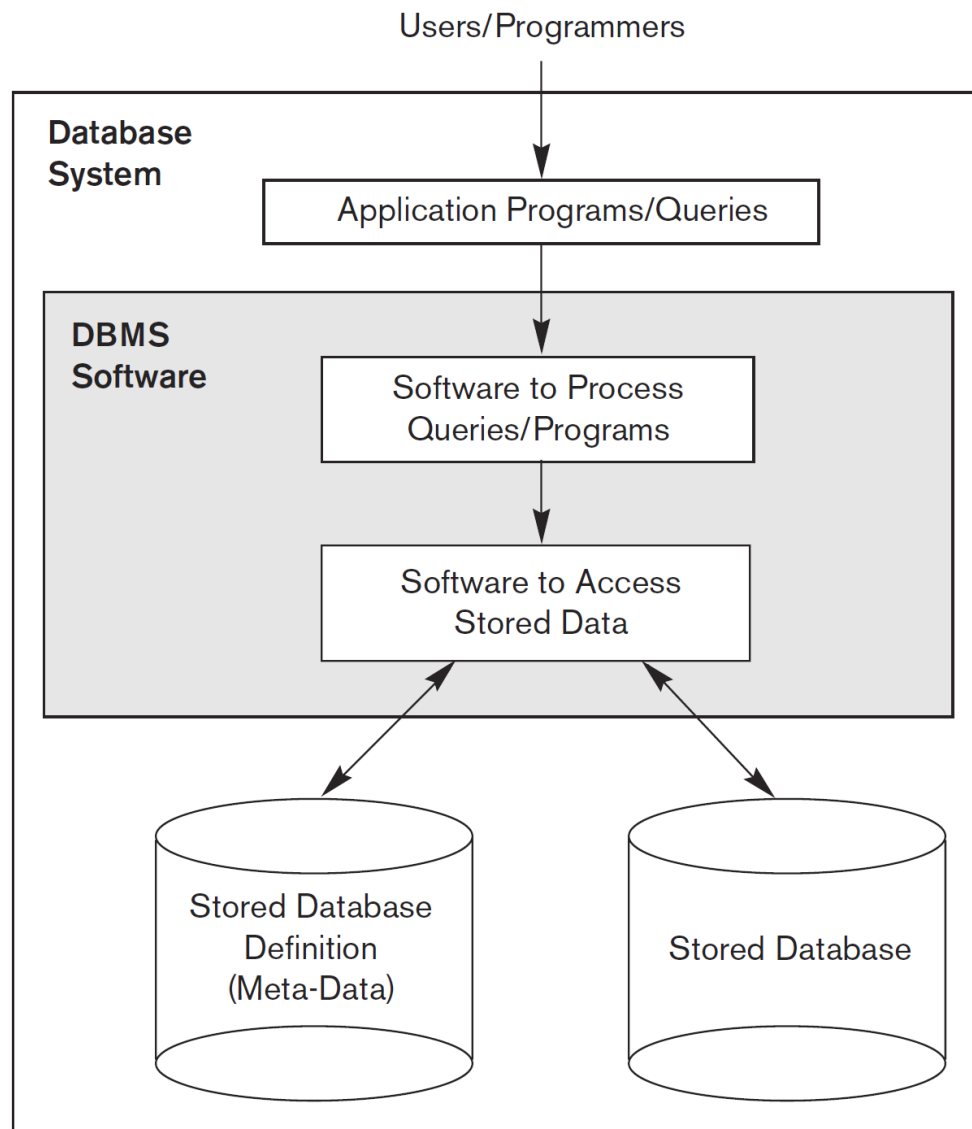



Relations, constraints and Joins

Definitions 1/2 - Recap

- Database Management System (DBMS) – A system that enables users to create and maintain a database.
- Database – A collection of related data. A collection of random data is not a database. Often the word database also, incorrectly, refers to a DBMS.
- Data – Known facts that can be recorded that has implicit meaning.
- Mini-World / Universe of Discourse – a database that represents some aspect of the real world. (Mostly here as the book refers to it)
- Meta data – Data about data, or data that gives context to other data.
- Schema – A definition of table structures and their relationships.
- SQL – Structured Query Language – A language to extract data from a database.

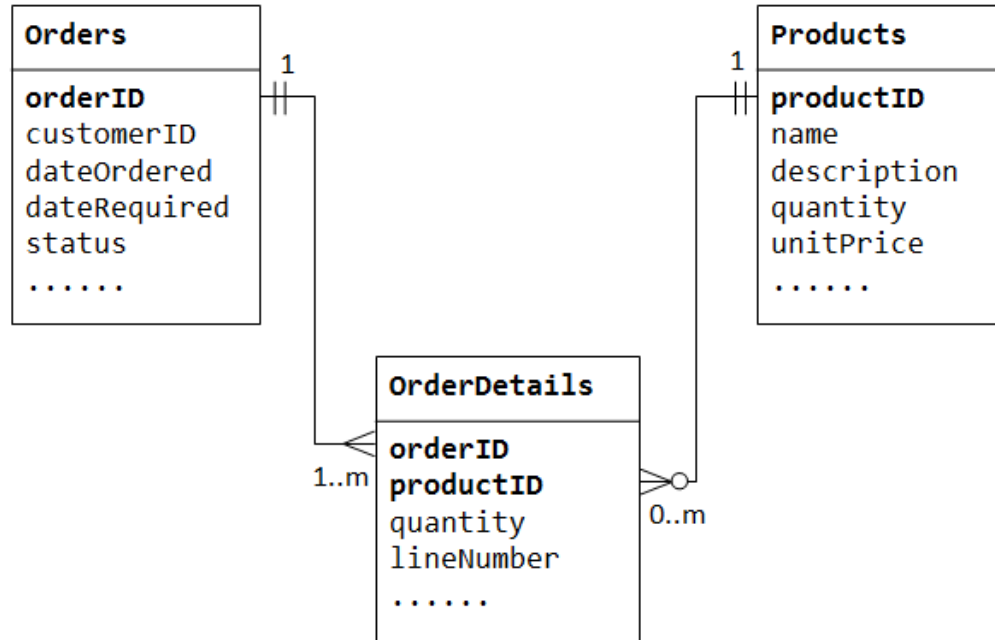


Definitions 2/2 - Recap



- Table – A collection of rows and columns that contains Cells.
- Row, Tuple, or Record – A collection of cells that together form a set of data that has a concrete Relation to each other. Shown as the horizontal line to the left.
- Column, Attribute, or Field – A collection of Cells that that are all the same type. They are a part of multiple Rows.
- Cell – A specific Rows Column. Ergo the intersection.
- Relation – The relationship between data in a Row.
- Value – The information saved in the specific Cell.
- View / Result Set – The table that is created in memory and sent to the client as a result of a query.

What is a relational database? - Recap



- Has a Schema that defines the Tables
- Uses multiple tables to store data
- Uses the notion of Primary Keys and Foreign keys to relate table content to each other.
- Uses SQL which makes CRUD (Create, Read, Update, Delete) operations on Schemas, Tables, and Rows.

Constraint types

- PRIMARY KEY – The unique identifier for the current row, which is always NOT NULL
- FOREIGN KEY – A key that refers to a primary key in a different table, signifying their relationship to each other
- UNIQUE – Two cells in this Column cannot be the same
- NOT NULL – This attribute HAS to be specified

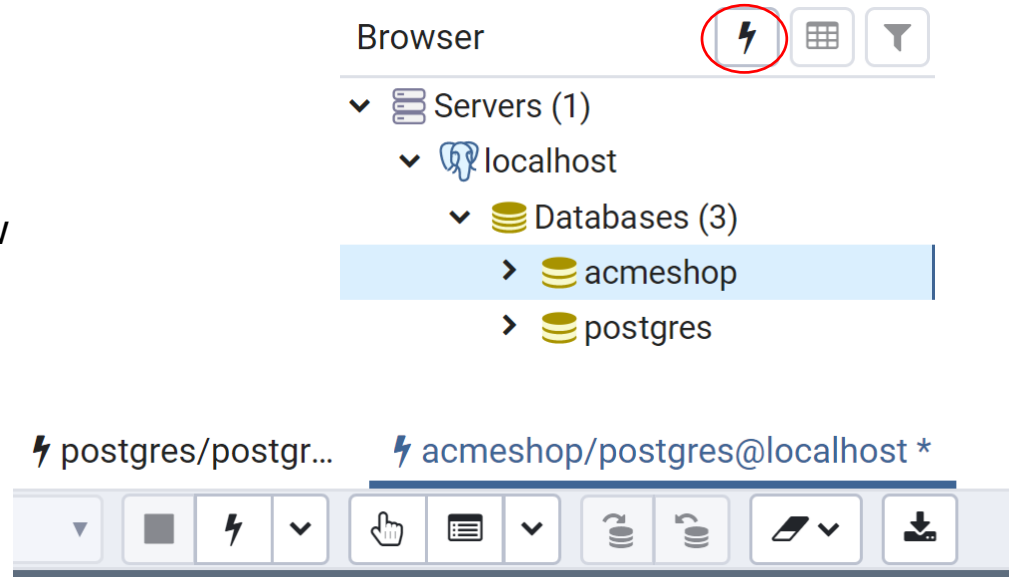
Constraints and Data types in SQL

```
CREATE TABLE account(  
    user_id serial PRIMARY KEY,  
    username VARCHAR (50) UNIQUE NOT NULL,  
    password VARCHAR (50) NOT NULL,  
    email VARCHAR (355) UNIQUE NOT NULL,  
    created_on TIMESTAMP NOT NULL,  
    last_login TIMESTAMP  
);
```

Creating a Database

- PostgreSQL only allows connections to one database at a time, and does not allow switching between them.
- This is not normal, and usually you can “use database” to attach to another database.
- Because of this restriction, run the command, pick the newly created database, and start a new SQL window.

```
-- creating a database  
create database AcmeShop;
```



Deleting a Database

```
drop database AcmeShop;
```

Active connections to a database will block a deletion attempt!

Tables and relationships

Customers Table

	id [PK] integer	username character varying (50)	password character varying (50)	email character varying (355)	created_on timestamp without time zone	last_login timestamp without time zone
1	1	John	myPassW0rd	john@acme.com	2020-02-13 10:40:41.660572	[null]
2	2	Anne	SomePassword	anne@acme.com	2020-02-13 10:40:41.660572	[null]

Products Table

	id [PK] integer	name character varying (150)	price real
1	1	Samsung Galaxy S20	7799.95
2	2	Samsung Galaxy S20 - Leathe...	799.95
3	3	iPhone 11 Pro	8899
4	4	iPhone 11 Pro - Leather Cover	399.5
5	5	Huawei P30 Lite	1664.5
6	6	Huawei P30 - Leather Cover	1664.5

Orders Table

	id [PK] integer	order_number character (10)	customer_id integer
1	1	DA-0001234	1
2	2	DA-0001235	1
3	3	DE-0001236	2
4	4	DE-0001237	2

Order_Lines Table

	id [PK] integer	order_id integer	product_id integer	amount integer
1	1	1	1	2
2	2	1	2	2
3	3	1	5	1
4	4	3	3	2
5	5	3	4	1
6	6	4	1	1

Creating Tables with Relationships

```
CREATE TABLE account(  
    id serial PRIMARY KEY,  
    username VARCHAR (50) UNIQUE NOT NULL,  
    password VARCHAR (50) NOT NULL,  
    email VARCHAR (355) UNIQUE NOT NULL,  
    created_on TIMESTAMP NOT NULL,  
    last_login TIMESTAMP  
);  
  
CREATE TABLE blog_entries(  
    id serial PRIMARY KEY,  
    header varchar(255) NOT NULL,  
    body TEXT NOT NULL,  
    created_by integer NOT NULL REFERENCES account(id)  
);
```

Constraint Enforcement

- ✓ `INSERT INTO account (username, password, email, created_on)`
`VALUES ('John', 'myPassW0rd', 'john@acme.com', NOW()); -- becomes user 1`
- ✓ `INSERT INTO blog_entries (header, body, created_by)`
`VALUES ('My article', 'my body text', 1); -- works!`
- ✓ `INSERT INTO blog_entries (header, body, created_by)`
`VALUES ('My article', 'my body text', 6); -- ERROR!`

```
ERROR: insert or update on table "blog_entries" violates foreign key constraint "blog_entries_created_by_fkey"  
DETAIL: Key (created_by)=(6) is not present in table "account".  
SQL state: 23503
```

Deleting tables, and tables with constraints

```
CREATE TABLE account(  
  id serial PRIMARY KEY,  
  username VARCHAR (50) UNIQUE NOT NULL,  
  password VARCHAR (50) NOT NULL,  
  email VARCHAR (355) UNIQUE NOT NULL,  
  created_on TIMESTAMP NOT NULL,  
  last_login TIMESTAMP  
);
```

```
CREATE TABLE blog_entries(  
  id serial PRIMARY KEY,  
  header varchar(255) NOT NULL,  
  body TEXT NOT NULL,  
  created_by integer NOT NULL REFERENCES account(id)  
);
```

```
drop table account;
```

→ Delete in order, where non of a tables attributes has a constraint from another table

```
ERROR: cannot drop table account because other objects depend on it  
DETAIL: constraint blog_entries_created_by_fkey on table blog_entries depends on table account  
HINT: Use DROP ... CASCADE to drop the dependent objects too.  
SQL state: 2BP01
```

Altering tables and deleting them

```
-- creating and altering tables
```

```
CREATE TABLE my_table(  
    id serial PRIMARY KEY,  
    my_attribute VARCHAR (50) UNIQUE NOT NULL  
);
```

```
ALTER TABLE my_table ADD COLUMN my_new_coulmn TIMESTAMP;
```

```
ALTER TABLE my_table ALTER COLUMN my_new_coulmn TYPE varchar(50);
```

```
ALTER TABLE my_table DROP COLUMN my_new_coulmn;
```

```
DROP TABLE my_table;
```

Inserts (CRUD)

```
INSERT INTO account (username, password, email, created_on)
VALUES ('John', 'myPassW0rd', 'john@acme.com', NOW());
```

```
INSERT INTO account (username, password, email, created_on)
VALUES ('Anne', 'myPassW0rd', 'anne@acme.com', NOW());
```








	user_id [PK] integer	username character varying (50)	password character varying (50)	email character varying (355)	created_on timestamp without time zone	last_login timestamp without time zone
1	4	John	myPassW0rd	john@acme.com	2020-02-06 11:43:44.158522	[null]
2	5	Anne	myPassW0rd	anne@acme.com	2020-02-06 11:43:44.158522	[null]

Selects (CRUD)

```
SELECT * FROM account;
```

```
SELECT username, created_on FROM account WHERE email = 'john@acme.com';
```

```
SELECT username, created_on FROM account WHERE email LIKE '%anne%';
```

	 user_id [PK] integer 	username character varying (50) 	password character varying (50) 	email character varying (355) 	created_on timestamp without time zone 	last_login timestamp without time zone 
1	1	John	myPassW0rd	john@acme.com	2020-02-06 11:41:11.729203	[null]

Updates (CRUD)

```
UPDATE account SET password = 'newPassW0rd' WHERE username = 'Anne';
```

	<div>user_id</div> <div>[PK] integer</div>	<div>username</div> <div>character varying (50)</div>	<div>password</div> <div>character varying (50)</div>	<div>email</div> <div>character varying (355)</div>	<div>created_on</div> <div>timestamp without time zone</div>	<div>last_login</div> <div>timestamp without time zone</div>
1	1	John	myPassW0rd	john@acme.com	2020-02-06 11:41:11.729203	[null]
2	2	Anne	newPassW0rd	anne@acme.com	2020-02-06 11:42:13.27506	[null]

Delete (CRUD)

```
DELETE FROM account WHERE email = 'john@acme.com';
```

	<div>user_id</div> <div>[PK] integer</div>	<div>username</div> <div>character varying (50)</div>	<div>password</div> <div>character varying (50)</div>	<div>email</div> <div>character varying (355)</div>	<div>created_on</div> <div>timestamp without time zone</div>	<div>last_login</div> <div>timestamp without time zone</div>
1	5	Anne	myPassW0rd	anne@acme.com	2020-02-06 11:43:44.158522	[null]

BEWARE of doing this: `DELETE FROM account;`

Querying Your Result Set - Nested Queries

```
CREATE TABLE account(
  id serial PRIMARY KEY,
  username VARCHAR (50) UNIQUE NOT NULL,
  password VARCHAR (50) NOT NULL,
  email VARCHAR (355) UNIQUE NOT NULL,
  created_on TIMESTAMP NOT NULL,
  last_login TIMESTAMP
);

CREATE TABLE blog_entries(
  id serial PRIMARY KEY,
  header varchar(255) NOT NULL,
  body TEXT NOT NULL,
  created_by integer NOT NULL REFERENCES account(id)
);
```

```
select * from account, blog_entries
where blog_entries.created_by = account.id;
```

	id integer	username character varying (50)	password character varying (50)	email character varying (355)	created_on timestamp without time zone	last_login timestamp without time zone
1	1	John	myPassW0rd	john@acme.com	2020-02-13 11:55:37.125376	[null]
2	1	John	myPassW0rd	john@acme.com	2020-02-13 11:55:37.125376	[null]
3	1	John	myPassW0rd	john@acme.com	2020-02-13 11:55:37.125376	[null]
4	2	Anne	SomePassword	anne@acme.com	2020-02-13 11:55:42.35475	[null]
5	1	John	myPassW0rd	john@acme.com	2020-02-13 11:55:37.125376	[null]

```
select username, email, created_by from (
  select * from account, blog_entries
  where blog_entries.created_by = account.id
) as result_set where created_by = 2;
```

	username character varying (50)	email character varying (355)	created_by integer
1	Anne	anne@acme.com	2

Join types

→ INNER JOIN

→ For each row R1 of T1, the joined table has a row for each row in T2 that satisfies the join condition with R1.

→ LEFT OUTER JOIN

→ First, an inner join is performed. Then, for each row in T1 that does not satisfy the join condition with any row in T2, a joined row is added with null values in columns of T2. Thus, the joined table always has at least one row for each row in T1.

→ RIGHT OUTER JOIN

→ First, an inner join is performed. Then, for each row in T2 that does not satisfy the join condition with any row in T1, a joined row is added with null values in columns of T1. This is the converse of a left join: the result table will always have a row for each row in T2.

→ FULL OUTER JOIN

→ First, an inner join is performed. Then, for each row in T1 that does not satisfy the join condition with any row in T2, a joined row is added with null values in columns of T2. Also, for each row of T2 that does not satisfy the join condition with any row in T1, a joined row with null values in the columns of T1 is added.

→ Source: <https://www.postgresql.org/docs/9.2/queries-table-expressions.html>

Inner Join


```
SELECT * FROM t1 INNER JOIN t2 ON t1.num = t2.num;
```

```
SELECT * FROM t1,t2 WHERE t1.num = t2.num;
```

num integer	name character varying (10)	num integer	value character varying (10)
1	a	1	xxx
3	c	3	yyy

```
SELECT * FROM t1 INNER JOIN t2 USING (num);
```

num integer	name character varying (10)	value character varying (10)
1	a	xxx
3	c	yyy



T1 table

num integer	name character varying (10)
1	a
2	b
3	c

T2 table

num integer	value character varying (10)
1	xxx
3	yyy
5	zzz


Left Outer Join

```
SELECT * FROM t1 LEFT JOIN t2 ON t1.num = t2.num;
```

num integer	name character varying (10)	num integer	value character varying (10)
1	a	1	xxx
2	b	[null]	[null]
3	c	3	yyy

```
SELECT * FROM t1 LEFT JOIN t2 USING (num);
```

num integer	name character varying (10)	value character varying (10)
1	a	xxx
2	b	[null]
3	c	yyy



T1 table

num integer	name character varying (10)
1	a
2	b
3	c


T2 table

num integer	value character varying (10)
1	xxx
3	yyy
5	zzz

Right Outer Join

```
SELECT * FROM t1 RIGHT JOIN t2 ON t1.num = t2.num;
```

num integer	name character varying (10)	num integer	value character varying (10)
1	a	1	xxx
3	c	3	yyy
[null]	[null]	5	zzz



T1 table

num integer	name character varying (10)
1	a
2	b
3	c


T2 table

num integer	value character varying (10)
1	xxx
3	yyy
5	zzz

Full Outer Join

```
SELECT * FROM t1 FULL JOIN t2 ON t1.num = t2.num;
```

num integer	name character varying (10)	num integer	value character varying (10)
1	a	1	xxx
2	b	[null]	[null]
3	c	3	yyy
[null]	[null]	5	zzz



T1 table

num integer	name character varying (10)
1	a
2	b
3	c

T2 table

num integer	value character varying (10)
1	xxx
3	yyy
5	zzz

Views – Creating Virtual Tables

```
CREATE VIEW MyCustomView AS
```

```
    SELECT email, username FROM account, blog_entries  
    WHERE blog_entries.created_by = account.id;
```

} Query defining the view content

```
SELECT * FROM MyCustomView;
```

	email character varying (355)	username character varying (50)
1	john@acme.com	John
2	john@acme.com	John
3	john@acme.com	John
4	anne@acme.com	Anne
5	john@acme.com	John

- Can join and simplify multiple tables into a view.
- Result of the query in the view is dynamically updated, when new database content is added.
- Can hide complexity of data
- Takes very little space as only the query is stored

Live Demo

Pay attention, and don't try to replicate what I do right now!

You will have time to do that afterwards.

Exercises

- Make sure PostgreSQL works! Ask instructors for help! It will be close to impossible to complete the course without doing the exercises.
- Create a database using a query.
- Create the tables from the next page WITH constraints in the new database.
- Create the content for all of the tables.
- Create the following queries:
 - Add another column to the Product Table called Manufacturer as a VARCHAR(250) using the ALTER TABLE command.
 - Query all orders with the product names bought, and the amount bought for each product, as well as the order_number and customer email.
 - Explore the different types of joins, Inner, left outer, right outer and full outer.
- Create a view with the query from above.
- Query the new view but look for a specific order number.

Create these tables with the same values

Customers Table

	id [PK] integer	username character varying (50)	password character varying (50)	email character varying (355)	created_on timestamp without time zone	last_login timestamp without time zone
1	1	John	myPassW0rd	john@acme.com	2020-02-13 10:40:41.660572	[null]
2	2	Anne	SomePassword	anne@acme.com	2020-02-13 10:40:41.660572	[null]

Products Table

	id [PK] integer	name character varying (150)	price real
1	1	Samsung Galaxy S20	7799.95
2	2	Samsung Galaxy S20 - Leathe...	799.95
3	3	iPhone 11 Pro	8899
4	4	iPhone 11 Pro - Leather Cover	399.5
5	5	Huawei P30 Lite	1664.5
6	6	Huawei P30 - Leather Cover	1664.5

Orders Table

	id [PK] integer	order_number character (10)	customer_id integer
1	1	DA-0001234	1
2	2	DA-0001235	1
3	3	DE-0001236	2
4	4	DE-0001237	2

Order_Lines Table

	id [PK] integer	order_id integer	product_id integer	amount integer
1	1	1	1	2
2	2	1	2	2
3	3	1	5	1
4	4	3	3	2
5	5	3	4	1
6	6	4	1	1